Packaging and containerizing of bioinformatics software: advances, challenges, and opportunities

Unsorted list of authors

- **Sharon Waymost**
  [0000-0003-1176-5386](#) · [sbpw](#)
  Computer Science Department, Samueli School of Engineering, University of California, Los Angeles

- **Serghei Mangul**
  [0000-0003-4770-3443](#) · [smangul1](#)
  Department of Clinical Pharmacy, School of Pharmacy, University of Southern California

- **Neha Rajkumar**

  Department of Bioengineering, Samueli School of Engineering, University of California, Los Angeles

- **Ram Ayyala**

  Department of Neuroscience, School of Life Sciences, University of California Los Angeles

- **Heng Li**

  None

- **Nathan LaPierre**

  None

- **André Santos**

  None

- **Titus Brown**

  University of California, Davis

- **Casey S. Greene**
  [0000-0001-8713-9213](#)
  Department of Systems Pharmacology and Translational Therapeutics, Perelman School of Medicine, University of Pennsylvania; Childhood Cancer Data Lab

- **Brendan Lawlor**
  [0000-0002-2250-0669](#) · [blawlor](#) · [blawlor](#)
  Department of Computer Science, Cork Institute of Technology, Cork, Ireland.

- **Qiyang Hu**

  None

- **Jaqueline J. Brito**
  [ID] 0000-0002-7158-3253 · [GitHub] jaquejbrito · [Twitter] jaquejbrito
  Department of Clinical Pharmacy, School of Pharmacy, University of Southern California

- **Thiago Mosqueiro**

  None

Abstract

Introduction

The rapid advancement of omics and sequencing technologies has led to an accelerated growth of genomic data. Driven by this data, there are similar increases in the number of available bioinformatics software tools, often containing novel and diverse algorithms of particular interest to biomedical researchers.[1,2]. Biomedical researchers wanting to use this software often have access to high-performance clusters ("clusters"), but lack the operating system-level permissions and advanced skills required to install and run the software. None of the necessary skills are currently included in the traditional life science curriculum at major universities. Usability is further hindered by the absence of any standardization and the wide variety of software development tools employed. This is directly observable in many academic software tools lacking a user-friendly interface[5].

As the dependence of biomedical researchers on computational software continues to increase, software developers need to consider more user-friendly distribution and install methods. Modern software distribution, having already faced user-friendliness issues, is already becoming more reliant on platforms such as package managers[6,7] and containers[8]. Both promise to simplify software development while increasing the usability and reproducibility of biomedical research[8,9].

Package managers first appeared nearly thirty years ago as software developers sought to streamline the entire installation process. To install via package manager, the user must only specify the desired software, called a "package;" the download, installation, configuration, and dependencies are all handled by the package manager. Many operating systems have built-in package managers (e.g. APT), which may not be available to cluster users, while others must be downloaded and installed by the user. Some package managers (e.g. Conda[6]) are programming language agnostic, while others are designed for a particular language (e.g pip[10]).

A more recent software distribution solution is containerization. The end-user downloads a container "image" that includes the software, dependencies, and anything else necessary to run the software. Though the imaged software is not typically itself installed, many do require the installation of containerization runtime software. When run, the runtime software creates a consistent, isolated sandbox environment, then runs the imaged software inside the sandbox[11]. This sandbox design makes the images both highly portable (compatible with different computers) and easily shareable (transferable between different computers), which has already led to wide adaptation in bioinformatics[7,12].

In addition to the ease of installation and use offered by both package managers and containers, such platforms must also meet the biomedical community's need for compatibility with high performance clusters. However, the relative performance of these package managers and containers remains unknown. Our review summarizes developing practices across the most common package managers and containers, while discussing the challenges, advantages, and limitations of using them from both the developer and the user perspectives. By taking a survey of all the available package managers and container software, there is now a comprehensive list of the different attributes of each one, informing the community so that people can make educated decisions about which package manager or container makes the most sense for their project. We also propose principles that can make packaging and containerizing of bioinformatics software more sustainable and reproducible, ultimately increasing the usability of bioinformatics software.

Discussion

Existing problems with software distribution and installation

The installation process of bioinformatics research software is typically a multi-step process, starting with the end-user locating and downloading the software. Next is the actual installation, during which the end-user must determine what dependencies are missing and resolve them by installing the required software[5]. Even in cases where the end-user is familiar with this process, they are typically installing on high-performance clusters where they are constrained by user-level permissions that prevent them from following standard installation procedures.

- root access limitations
- reproducibility of findings
- version conflicts
- dependency resolution

Definitions and explanations of distribution system types

Package Managers

Package managers appeared nearly thirty years ago as software developers sought to automate and standardize the software installation process. The earliest package managers (e.g. APT and RPM) were built into Linux operating systems and lacked now common features like dependency resolution. Modern package managers are available for all major operating systems, with both integrated and user-installed options available. Some are package managers are limited to software written in a specific programming language (e.g. pip distributes software written in Perl), while others are limited to sofware with a particular purpose (e.g. Bioconda distributes omics software).

Every package managers uses its own package format, which follows the general structure shown in Figure 1.
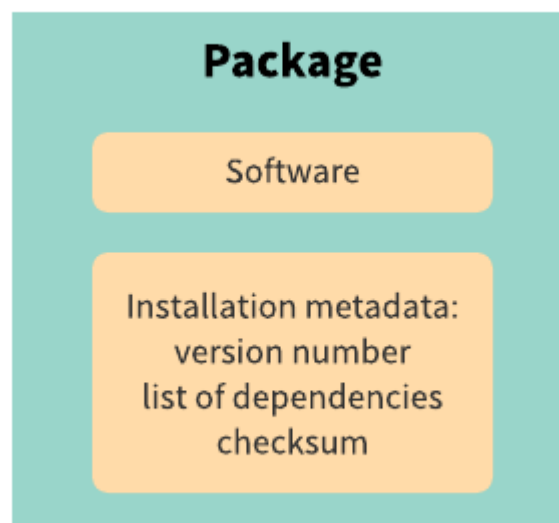
Figure 1: Schematic showing standard package contents. Though the exact contents varies between different package formats, all include the software and relevant metadata necessary to complete installation.

The metadata gives the package manager the necessary information to install the packaged software.

To use a package manager, a user instructs the package manager to install software, then the package manager performs all the necessary steps to install the software as shown in Figure 2. Assuming the software is not already installed, the package manager retrieves the appropriate package from the repository, a server that stores all of the available packages. The package manager then uses the package's metadata to determine what it needs to do to complete installation. For most packages, this includes verifying the installation of all dependencies, which are other pieces of software that must be installed prior to the initial software being installed.
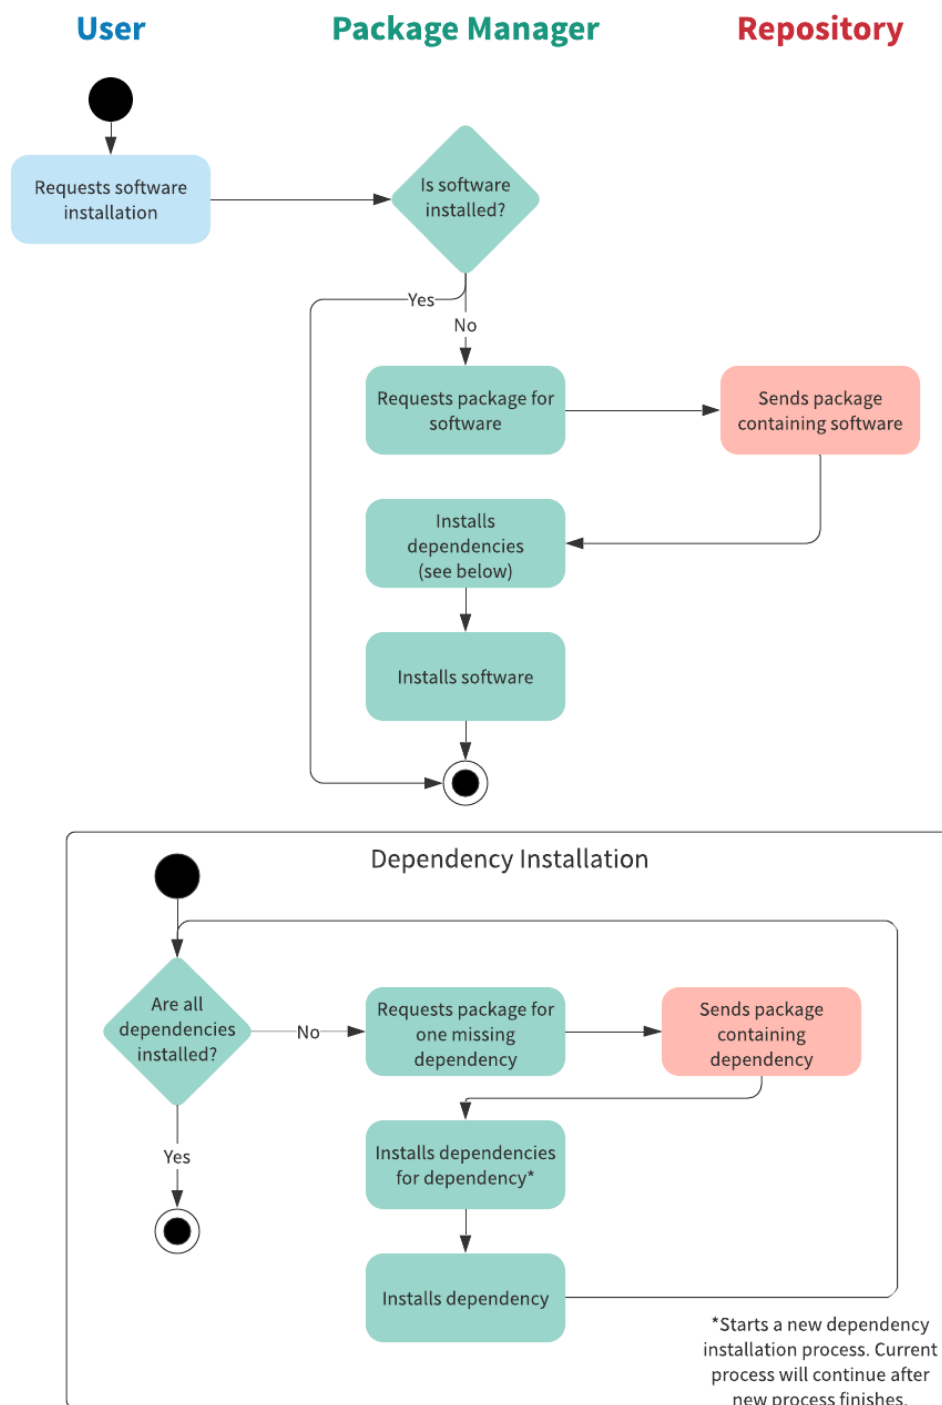
Figure 2: Standard workflow for installing software via a package manager. The user, usually an admin, asks the package manager to install a specific piece of software. If the software is not already installed, the package manager fetches the approrpriate package from the repository. If any of the dependencies are not already installed, the package manager "loops" through the missing dependencies. Once all dependencies are installed, the original software is installed.

The package manager verifies each dependency by confirming that the dependency is already installed or installing the dependency if it is missing. For missing dependencies, the package manager retrieves the dependency's package from the repository and starts the installation procedure for that package. Once all the dependencies are installed, the initially requested software is installed. However, the package manager often goes through several iterations of this process because every dependency can have it's own list of dependencies, in which case each of the dependency's dependencies must be verified through the same process.

In most cases, a package manager is able to install all of the missing dependencies, but there are two relatively common problems that a package manager cannot overcome. The first is circular dependencies, which occurs when software A has a dependency on software B and software B has a dependency on software A. A package manager cannot install A without installing B first, but B cannot be installed without installing A first. The second is conflicting versions, which occurs when software A and software B have the same dependency, but require different versions. Package managers only allow one version of software to be installed at a time, preventing both A and B from having their depedencies met simultaneously.

Administrators often block users from installing software to prevent unwanted changes to both the operating system and existing sofware. This includes blocking access to package managers built into the operating system, as well as preventing users from installing their own.

- benefits for the developer
  - mature technology - higher degree of familiarity
  - allows dependency specification (including versions) -limitations for the developer
  - can't always use to install missing dependencies for end-user
- benefits for the end-user
  - package size is minimal (dependencies aren't duplicated)
  - installs missing dependencies
- limitations for the end-user
  - not always accessible (unless admin user)
  - can't install multiple versions of same software
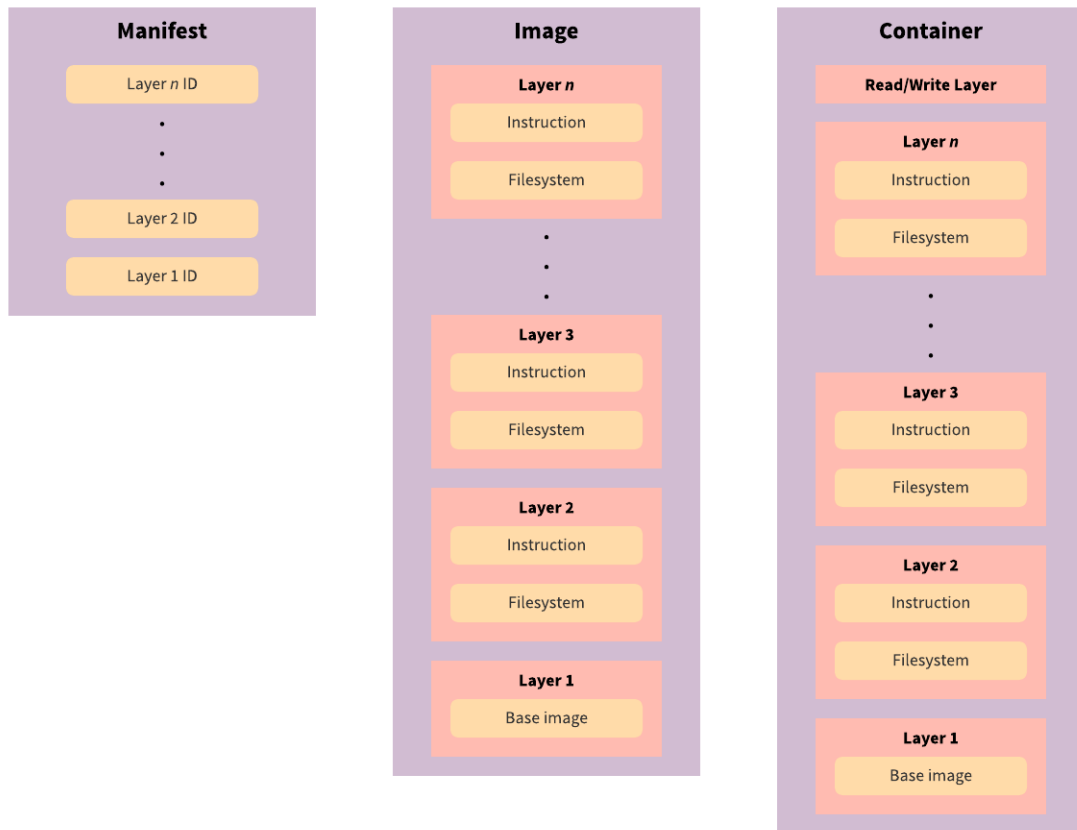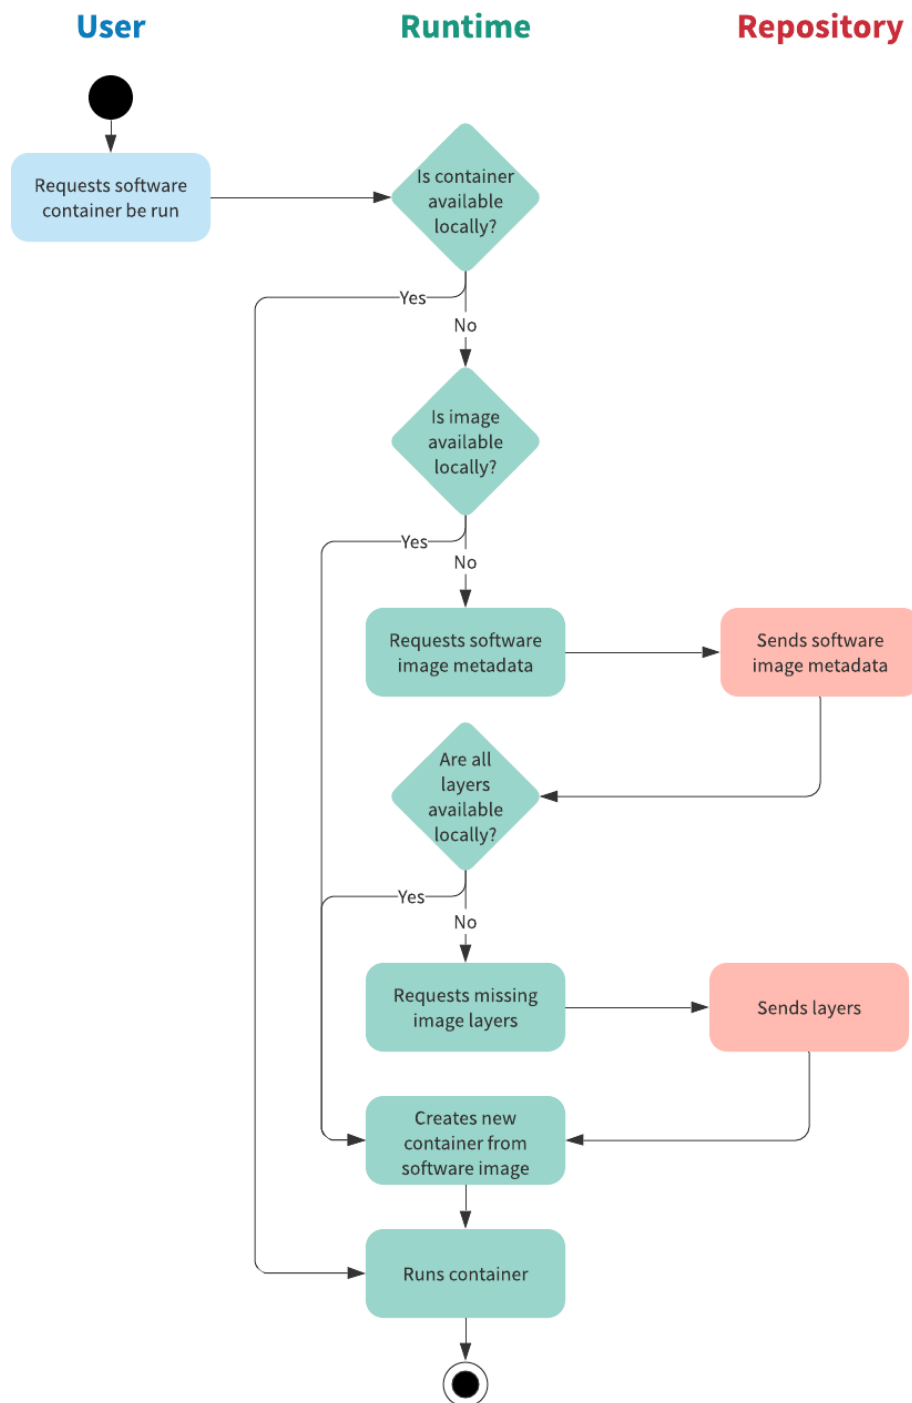- containerization
  - definition

Figure 3: Schematic showing standard containerization objects.

Figure 4: Running containerized software

- benefits for the developer
  - include specific versions of dependencies
  - known running environment
    - fewer test variables
    - reproducibility of results
- limitations for the developer
  - learn a new system instead of focusing on research
- benefits for the end-user
  - no installation (except possible runtime)
  - no dependency issues
  - sandbox provides computer system security
- limitations for the end-user

- container size
    - duplication of dependencies
    - root access requirement to install runtime
    - configuration in cluster
- centralized repositories
    - definition
        - channels
    - benefits
        - known download site
        - hosting is taken of
   - limitations
        - repo specific restrictions



Figure 5: Relative popularity of package managers.

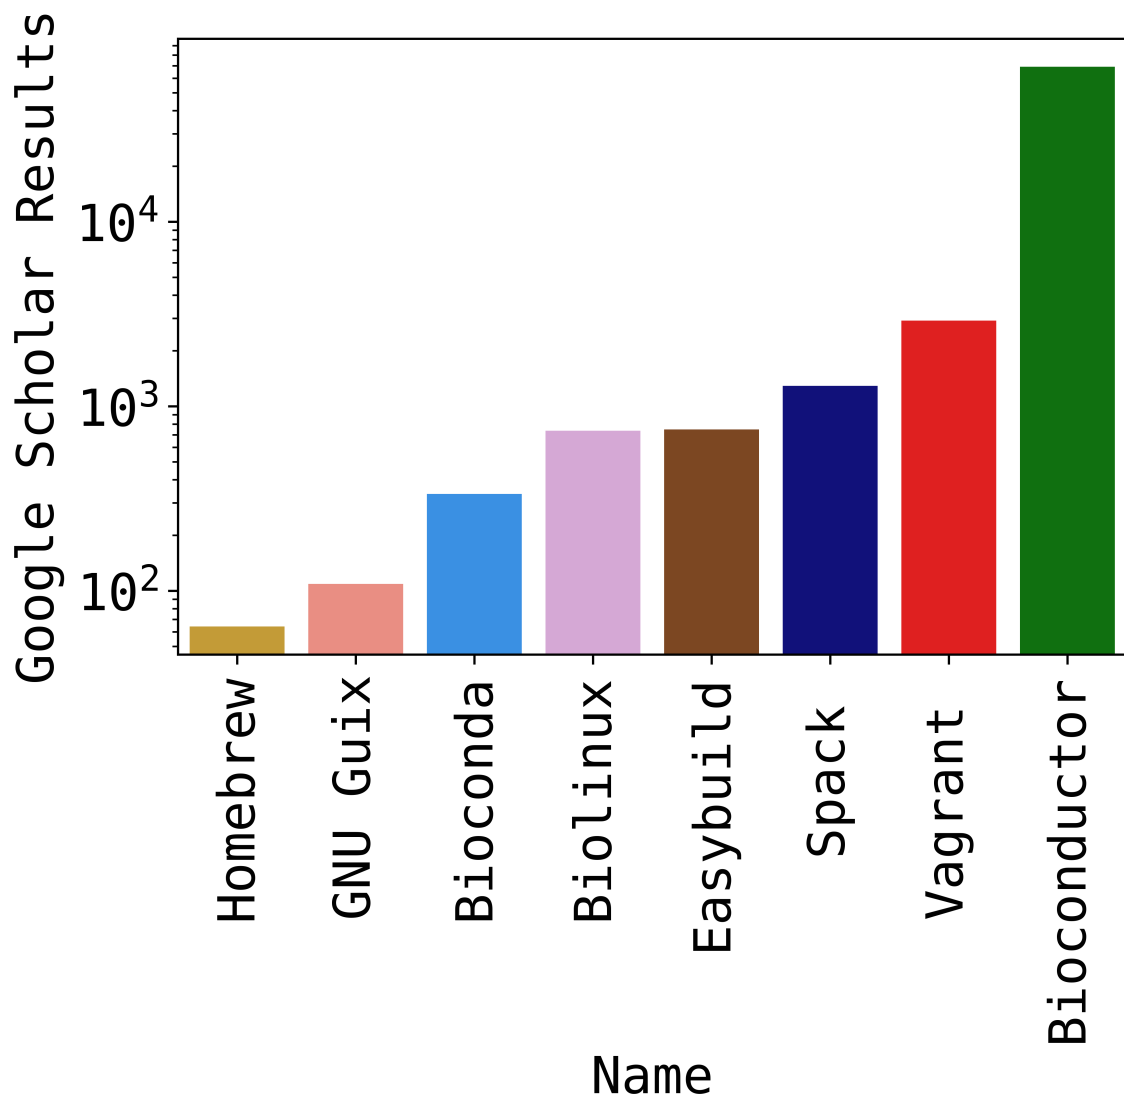**Figure 6:** Usage of package managers in academia based on citations.

**Figure 7: Usage of package managers in academia based on Google Scholar Results.**

**Figure 8: Total number of packages.**

Figure 9: Packages heatmap

**Figure 10: Overlap in available packages between omics package managers.**

Figure 11: Availablity of Bioconda's packages in non-omics package managers.

Figure 12: Availablity of Bioinformatics Tools' packages in non-omics package managers.

Figure 13: Availability of omicX's packages in non-omics package managers.

Methods

5 Google Trends [1]/ was used to assess the relative average popularity of each package manager. Package managers were omitted when Google Trends did not have enough data to return a result. The method detailed in "How do you compare large number of items in Google Trends?" [2]/ was used to overcome Google Trends' limitation of five search terms.

6 For package managers introduced with a major publication, the number of citations for their publication was found via Google Scholar [3]/ and recorded. Package managers without a major publication were excluded.

7 Each package manager's name was used as a search term on Google Scholar [3]/ and the number of results was recorded. Package manager names that returned results for things other than the package manager were omitted.

8 The initial list of all available sofware for each package manager was compiled by parsing its online repository with WebScraper [4]/ (AppImage [5], Bioconda [6], Bioconductor [7], BioContainers [**???**], CRAN [8], Docker [9], EasyBuild [10], Flatpak [11], GNU Guix [12], omicX [13], pip [14], Singularity [15], Snap [16], Spack [17], Vagrant [18], and Zero Install [19]) or by querying the package manager directly (APT, Biocontainers, Conda, Homebrew). Each list converted to a set with pandas [20]/ and NumPy [21]/, which removed all capitalization and duplicates. The number of entries in a package manager's set was taken to be the total number of packages available for that package manager.

9 The sets obtained in [8] were compared with the set intersection function to find the number of commonalities between all pairs of package managers.

?? The sets obtained for Bioconda, Bioconductor, and omicX in [8] were compared with the set intersection function to find the number of commonalities between all combinations of the three package managers.

11 The sets obtained for Bioinformatics Tools and the non-omics package managers in [8] were compared with the set intersection function to find the number of commonalities between Bioconda and the non-omics package managers.

12 The sets obtained for Bioinformatics Tools and the non-omics package managers in [8] were compared with the set intersection function to find the number of commonalities between Bioinformatics Tools and the non-omics package managers.

13 The sets obtained for omicX and the non-omics package managers in [8] were compared with the set intersection function to find the number of commonalities between omicX and the non-omics package managers.

Glossary

Acknowledgements

Author Contributions

References

1. **Google Trends**
   Google Trends
   https://trends.google.com/trends

2. **How do you compare large numbers of items in Google Trends?**
   jonathanbriggs
   Digital Jobs To Be Done (2017-07-10) https://digitaljobstobedone.com/2017/07/10/how-do-you-compare-large-numbers-of-items-in-google-trends/

3. **Google Scholar** https://scholar.google.com/

4. **Web Scraper - The #1 web scraping extension** https://webscraper.io/

5. **Apps in AppImage format**
   AppImageHub
   AppImageHub https://appimage.github.io/apps/

6. https://bioconda.github.io/conda-recipe_index.html

7. **Bioconductor - BiocViews** http://bioconductor.org/packages/release/BiocViews.html#___Software

8. **CRAN Packages By Name** https://cran.r-project.org/web/packages/available_packages_by_name.html

9. **Docker Hub** https://hub.docker.com/search?q=&type=image&page=1

10. **List of supported software — EasyBuild v4.2.2 documentation (release 20200708.0)**
    https://easybuild.readthedocs.io/en/latest/version-specific/Supported_software.html

11. **Flathub—An app store and build service for Linux** https://flathub.org/apps/category/All

12. **Packages — GNU Guix** https://guix.gnu.org/packages/

13. https://omictools.com/software?page=1

14. **Simple index** https://pypi.org/simple

15. **singularity-hub**
    vsoch
    Singularity Hub https://singularity-hub.org/apps

16. **Install Linux apps using the Snap Store | Snapcraft**
    @snapcraftio
    Snapcraft https://snapcraft.io/store

17. **Package List — Spack 0.15.4 documentation** (2020-08-28)
    https://spack.readthedocs.io/en/latest/package_list.html

18. **Vagrant Cloud by HashiCorp**
    Vagrant Cloud by HashiCorp
     https://app.vagrantup.com

19. http://roscidus.com/0mirror/feed-list

20. **pandas - Python Data Analysis Library** https://pandas.pydata.org/

21. **NumPy** https://numpy.org/

Tables

Distribution System Name	URL	Publication	Type	License
AppImage	https://appimage.org	-	containerization	MIT
APT	https://wiki.debian.org/Apt	-	package manager	GNU GPL 2+
Bioconda	https://bioconda.github.io	Grüning et al, 2018	package manager	MIT
Bioconductor	https://www.bioconductor.org	Gentleman et al, 2004	package manager	MIT
conda	https://docs.conda.io/en/latest/	-	package manager	3-Clause BSD
CRAN	https://cran.r-project.org/index.html	-	package manager	GNU GPL
Docker	https://www.docker.com	-	containerization	Apache 2.0
Easybuild	https://easybuilders.github.io/easybuild/	Hoste et al, 2012	package manager	GNU GPL 2
Flatpak	https://flatpak.org	-	containerization	LGPL
GNU Guix	https://guix.gnu.org	Courtès, 2013	package manager	GNU AGPL
Homebrew	https://brew.sh	-	package manager	2-Clause BSD
pip	https://pypi.org/project/pip/	-	package manager	MIT
Singularity	https://sylabs.io	-	containerization	3-Clause BSD
Snap	https://snapcraft.io	-	containerization	propriertary
Spack	https://spack.io	Gamblin et al, 2015	package manager	MIT or Apache
Vagrant	https://www.vagrantup.com	-	virtual machine	MIT
yum	http://yum.baseurl.org	-	package manager	
Zero Install	https://0install.net	-	package manager	GNU LGPL 2.1+

Distribution System Name	Supported Operating Systems	Supported Languages	Root to Install	Root to Run
AppImage	Linux	any	n/a	no
APT	Debian, Ubuntu	any	yes	yes
Bioconda	Linux, macOS, Windows	any	no	no
Bioconductor	Linux, macOS, Windows	R	no	no
conda	Linux, macOS, Windows	any	no	no

Distribution System Name	Supported Operating Systems	Supported Languages	Root to Install	Root to Run
CRAN	Linux, macOS, Windows	R	no	no
Docker	Linux, macOS, Windows	any	yes	no
Easybuild	Linux	any	no	no
Flatpak	Linux	any	no	no
GNU Guix	Linux	any	no	no
Homebrew	Linux, macOS	any	no	no
pip	Linux, macOS, Windows	Python	no	no
Singularity	Linux, macOS	any	yes	no
Snap	Linux	any	yes	no
Spack	Linux, macOS	any	no	no
Vagrant	Linux, macOS, Windows	any	yes	
yum	Linux, macOS, Windows	any	no	yes
Zero Install	Linux, macOS, Windows	any	no	no

Distribution System Name	First Release	Latest Release	Age in Years	Number of Releases	Number of Tools	Number of Bio Tools
AppImage	2014-01-24	2020-06-01	7	121		
APT	1998-03-31	2020-05-08	22	362		
Bioconda	2014-01-24	2016-09-06	7	39		
Bioconductor	2002-05-01	2020-04-28	17	37		
conda	2014-01-24	2020-04-13	6	261		
CRAN	1997-04-23	2020-02-29	22	29		
Docker*	2013-03-23	2020-06-01	7	121		
Easybuild	2012-11-09	2020-04-14	7	51		
Flatpak	2015-03-23	2020-04-03	5	128		
GNU Guix	2012-07-07	2020-04-15	7	23		
Homebrew	2009-05-20	2020-05-04	10	155		
pip	2009-01-20	2020-04-28	11	81		
Singularity	2012-07-07	2020-04-15	7	23		
snapd	2014-12-09	2020-07-15	5	232		
Spack	2014-07-09	2020-04-15	5	27		
Vagrant						
yum**	2002-06-08	2011-06-28	18	221		
Zero Install	2005-02-04	2020-05-04	15	145		

*Docker Engine

**need to find someone with a redhat license who can confirm numbers

Distribution System Name	Official Repository Name	Repository URL
AppImage	AppImageHub	https://appimage.github.io/apps/
APT	-	-
Bioconda	bioconda channel	https://github.com/bioconda/bioconda-recipes
Bioconductor	-	https://www.bioconductor.org/packages/release/BiocViews.html#___Software
conda	-	https://repo.anaconda.com/pkgs/
CRAN	-	https://cran.r-project.org/web/packages/available_packages_by_name.html
Docker	Docker Hub	https://hub.docker.com
Easybuild		
Flatpak	Flathub	https://flathub.org
GNU Guix	-	https://guix.gnu.org/packages/
Homebrew	Homebrew Formulae	https://formulae.brew.sh
pip	Python Package Index (PyPI)	https://pypi.org
Singularity	Singularity Hub	https://singularity-hub.org
Snap	Snapcraft	https://snapcraft.io/store
Spack	-	-
Vagrant	Vagrant Cloud	https://app.vagrantup.com/boxes/search
yum	-	-
Zero Install	-	https://apps.0install.net